

Objective:

Objective of this model is to predict whether the customer is prone to churn from the existing network to new network or not.

Project Plan :

Analyzing and Importing data:

Importing data and analyzing it to extract insights that supports for decision making process (like obtaining dependent and independent variables).

Data Cleansing:

- Check nulls (If null: Delete the row or replace with mean, median or mode)
- Foreign values
- Check voids
- Wrong format data

Minimize the Dimension of Dataframe:

- Check correlation between columns -Pearson correlation is used to check correlation between two continuous column.
- Highly correlated columns are obtained and one of them are dropped.

Encoding Data:

Categorical data is encoded (using label or one hot encoding based on the requirement) so that they're easily readable by the machines.

Model building:

- Splitting the data into train and test data sets
- Train the model using decision tree classifier
- Fit the data into the model
- Predicting the response of the test dataset
- Evaluating accuracy of the model

Optimizing the Model performance:

- Gini and Entropy algorithms are used to optimize the performance of the model developed.
- Maximum depth is given for the decision tree to avoid overfitting and underfitting

Importing libraries

```
import pandas as pd #Data manipulation
import numpy as np #Data manipulation
import matplotlib.pyplot as plt # Visualization
import seaborn as sns #Visualization
import plotly.express as px #Visualization
import plotly.graph_objs as go #Visualization
import os
```

```
## for statistical tests
```

```

import scipy
import statsmodels.formula.api as smf
import statsmodels.api as sm

## for machine learning
from sklearn import model_selection, preprocessing, feature_selection,
ensemble, linear_model, metrics, decomposition

## for explainer
#from lime import lime_tabular

```

```

pd.options.plotting.backend = "plotly"
plt.rcParams['figure.figsize'] = [8,5]
plt.rcParams['font.size'] =14
plt.rcParams['font.weight']= 'bold'
plt.style.use('seaborn-whitegrid')

```

```

df=pd.read_csv('Churn.csv')
df.head()

```

	c_AreaCode	c_InternationalPlan	c_Phone	c_State	c_VMailPlan	\
0	415	no	382-4657	KS	yes	
1	415	no	371-7191	OH	yes	
2	415	no	358-1921	NJ	no	
3	408	yes	375-9999	OH	no	
4	415	yes	330-6626	OK	no	

	q_AccountLength	q_CustServCalls	q_DayCalls	q_DayCharge
q_DayMins ... \				
0	128	1	110	45.07
265.1 ...				
1	107	1	123	27.47
161.6 ...				
2	137	0	114	41.38
243.4 ...				
3	84	2	71	50.90
299.4 ...				
4	75	3	113	28.34
166.7 ...				

	q_EveCharge	q_EveMins	q_InternationalCharge	q_InternationalMins
\				
0	16.78	197.4	2.70	10.0
1	16.62	195.5	3.70	13.7
2	10.30	121.2	3.29	12.2
3	5.26	61.9	1.78	6.6

4	12.61	148.3	2.73	10.1
---	-------	-------	------	------

	q_Internationalcalls	q_NightCalls	q_NightCharge	q_NightMins	\
0	3	91	11.01	244.7	
1	3	103	11.45	254.4	
2	5	104	7.32	162.6	
3	7	89	8.86	196.9	
4	3	121	8.41	186.9	

	q_VMailMessage	y_Churn
0	25	False.
1	26	False.
2	0	False.
3	0	False.
4	0	False.

[5 rows x 21 columns]

Data Cleansing

```
df.isnull().sum()
```

```
c_AreaCode      0
c_InternationalPlan  0
c_Phone         0
c_State         0
c_VMailPlan     0
q_AccountLength  0
q_CustServCalls  0
q_DayCalls      0
q_DayCharge     0
q_DayMins       0
q_EveCalls      0
q_EveCharge     0
q_EveMins       0
q_InternationalCharge  0
q_InternationalMins  0
q_Internationalcalls  0
q_NightCalls    0
q_NightCharge   0
q_NightMins     0
q_VMailMessage  0
y_Churn         0
dtype: int64
```

View summary of data

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 4617 entries, 0 to 4616
```

```
Data columns (total 21 columns):
```

#	Column	Non-Null Count	Dtype
0	c_AreaCode	4617 non-null	int64
1	c_InternationalPlan	4617 non-null	object
2	c_Phone	4617 non-null	object
3	c_State	4617 non-null	object
4	c_VMailPlan	4617 non-null	object
5	q_AccountLength	4617 non-null	int64
6	q_CustServCalls	4617 non-null	int64
7	q_DayCalls	4617 non-null	int64
8	q_DayCharge	4617 non-null	float64
9	q_DayMins	4617 non-null	float64
10	q_EveCalls	4617 non-null	int64
11	q_EveCharge	4617 non-null	float64
12	q_EveMins	4617 non-null	float64
13	q_InternationalCharge	4617 non-null	float64
14	q_InternationalMins	4617 non-null	float64
15	q_Internationalcalls	4617 non-null	int64
16	q_NightCalls	4617 non-null	int64
17	q_NightCharge	4617 non-null	float64
18	q_NightMins	4617 non-null	float64
19	q_VMailMessage	4617 non-null	int64
20	y_Churn	4617 non-null	object

```
dtypes: float64(8), int64(8), object(5)
```

```
memory usage: 757.6+ KB
```

```
v=df.groupby(['c_InternationalPlan'])['q_InternationalCharge'].mean()  
v
```

```
c_InternationalPlan
```

```
no      2.766931
```

```
yes      2.860045
```

```
Name: q_InternationalCharge, dtype: float64
```

```
#Visualization
```

```
import matplotlib.pyplot as plt
```

```
x=['Subscribed','Unsubscribed']
```

```
y=[2.77,2.86]
```

```
a=plt.bar(x,y,width=0.4)
```

```
plt.xlabel('Subscription')
```

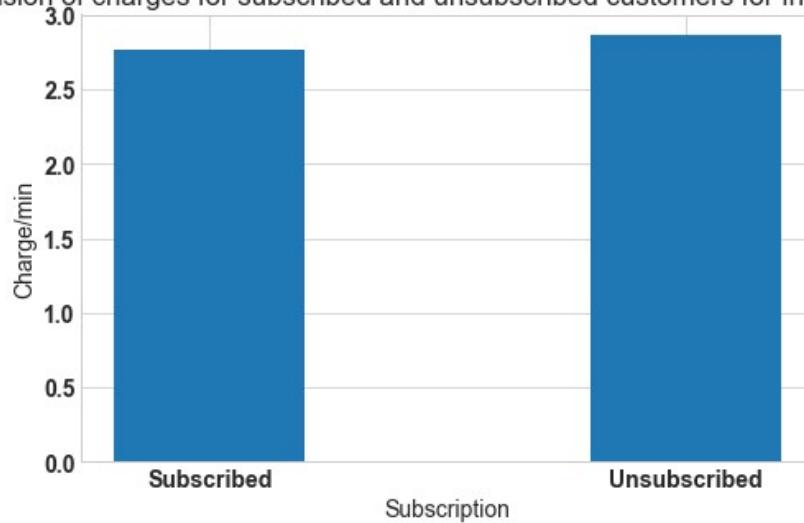
```
plt.ylabel('Charge/min')
```

```
plt.title('Comparision of charges for subscribed and unsubscribed  
customers for International calls')
```

```
a
```

```
<BarContainer object of 2 artists>
```

Comparison of charges for subscribed and unsubscribed customers for International calls



Inference :

From the plot it is observed that the charges for subscribed customer and unsubscribed customer varies for international calls.

```
df['DC/min']=df['q_DayCharge']/df['q_DayMins']
df['EC/min']=df['q_EveCharge']/df['q_EveMins']
df['NC/min']=df['q_NightCharge']/df['q_NightMins']
df['IC/min']=df['q_InternationalCharge']/df['q_InternationalMins']
df
```

	c_AreaCode	c_InternationalPlan	c_Phone	c_State	c_VMailPlan	\
0	415	no	382-4657	KS	yes	
1	415	no	371-7191	OH	yes	
2	415	no	358-1921	NJ	no	
3	408	yes	375-9999	OH	no	
4	415	yes	330-6626	OK	no	
...	
4612	510	no	345-7512	NY	yes	
4613	408	no	343-6820	NM	yes	
4614	408	no	338-4794	VT	yes	
4615	415	no	355-8388	MI	yes	
4616	415	no	409-6884	IN	no	

	q_AccountLength	q_CustServCalls	q_DayCalls	q_DayCharge
q_DayMins \				
0	128	1	110	45.07
265.1				
1	107	1	123	27.47
161.6				
2	137	0	114	41.38
243.4				
3	84	2	71	50.90

299.4				
4	75	3	113	28.34
166.7				
...
...				
4612	57	3	81	24.48
144.0				
4613	177	3	91	32.13
189.0				
4614	67	1	126	21.68
127.5				
4615	98	0	98	28.71
168.9				
4616	140	2	100	34.80
204.7				

	...	q_Internationalcalls	q_NightCalls	q_NightCharge
q_NightMins \				
0	...	3	91	11.01
244.7				
1	...	3	103	11.45
254.4				
2	...	5	104	7.32
162.6				
3	...	7	89	8.86
196.9				
4	...	3	121	8.41
186.9				
...
...				
4612	...	6	122	7.14
158.6				
4613	...	1	116	7.36
163.6				
4614	...	3	91	9.04
200.9				
4615	...	3	96	7.45
165.5				
4616	...	4	115	9.13
202.8				


	q_VMailMessage	y_Churn	DC/min	EC/min	NC/min	IC/min
0	25	False.	0.170011	0.085005	0.044994	0.270000
1	26	False.	0.169988	0.085013	0.045008	0.270073
2	0	False.	0.170008	0.084983	0.045018	0.269672
3	0	False.	0.170007	0.084976	0.044997	0.269697

4	0	False.	0.170006	0.085030	0.044997	0.270297
...
4612	25	False.	0.170000	0.084989	0.045019	0.270588
4613	29	False.	0.170000	0.084988	0.044988	0.270064
4614	33	False.	0.170039	0.085005	0.044998	0.270000
4615	23	False.	0.169982	0.085020	0.045015	0.269930
4616	0	False.	0.170005	0.085016	0.045020	0.270248

[4617 rows x 25 columns]

```
import matplotlib.pyplot as plt
fig=plt.figure()
fig.set_size_inches(25,5)
x=df['c_State']
y1=df['DC/min']
y2=df['EC/min']
y3=df['NC/min']
y4=df['IC/min']
a=plt.scatter(x,y1)
b=plt.scatter(x,y2)
c=plt.scatter(x,y3)
d=plt.scatter(x,y4)
plt.xlabel('States')
plt.ylabel('Charge/min')
plt.title('Comparision of charges/min for diff.time according to
states')
a
```


<matplotlib.collections.PathCollection at 0x1ffc2aead30>



Inference :

From the above plot it is observed that the charges doesn't vary much from state to state.
So the c_State column can be dropped

```
plt.figure(figsize=(20,5))  
ax = sns.countplot(x='c_State', hue="y_Churn", data=df)
```



```
max_acl=df['q_AccountLength'].max()  
min_acl=df['q_AccountLength'].min()  
print("The maximum account length :",max_acl)  
print("The minimum account length :",min_acl)
```

The maximum account length : 243

The minimum account length : 1

```
df.iloc[817]
```

c_AreaCode	510
c_InternationalPlan	no
c_Phone	355-9360
c_State	UT
c_VMailPlan	no
q_AccountLength	243
q_CustServCalls	2
q_DayCalls	92
q_DayCharge	16.24
q_DayMins	95.5
q_EveCalls	63
q_EveCharge	13.91
q_EveMins	163.7
q_InternationalCharge	1.78
q_InternationalMins	6.6
q_Internationalcalls	6
q_NightCalls	118
q_NightCharge	11.89
q_NightMins	264.2
q_VMailMessage	0
y_Churn	False.
DC/min	0.170052
EC/min	0.084973
NC/min	0.045004



```
IC/min          0.269697
Name: 817, dtype: object
```

```
df.iloc[245]
```

```
c_AreaCode          408
c_InternationalPlan no
c_Phone             331-5138
c_State             SC
c_VMailPlan         no
q_AccountLength     22
q_CustServCalls     0
q_DayCalls          107
q_DayCharge         18.75
q_DayMins           110.3
q_EveCalls          93
q_EveCharge         14.15
q_EveMins           166.5
q_InternationalCharge 2.57
q_InternationalMins  9.5
q_Internationalcalls 5
q_NightCalls        96
q_NightCharge        9.1
q_NightMins         202.3
q_VMailMessage      0
y_Churn             False.
DC/min              0.169991
EC/min              0.084985
NC/min              0.044983
IC/min              0.270526
Name: 245, dtype: object
```

```
import matplotlib.pyplot as plt
fig=plt.figure()
fig.set_size_inches(15,5)
x1=['DC(Old)', 'DC(New)']
y1=[0.169,0.170]
x2=['EC(Old)', 'EC(New)']
y2=[0.084,0.085]
x3=['NC(Old)', 'NC(New)']
y3=[0.044,0.045]
x4=['IC(Old)', 'IC(New)']
y4=[0.26,0.27]
plt.bar(x1,y1,width=0.7)
plt.bar(x2,y2,width=0.7)
plt.bar(x3,y3,width=0.7)
plt.bar(x4,y4,width=0.7)
plt.xlabel('Account Length')
plt.ylabel('Charge/min')
plt.title('Comparision of charges for new customer and existing
```

```
customer')
plt.show()
```



Inference:

The charges for low account length customer and high account length customer remains the same. So account length column can be dropped.

```
statechurncount=df.groupby("c_State").y_Churn .value_counts()
statechurncount
```

```
c_State  y_Churn
AK      False.    64
        True.      5
AL      False.   104
        True.     10
AR      False.    62
...
WI      True.      8
WV      False.   128
        True.     21
WY      False.    96
        True.     10
Name: y_Churn, Length: 102, dtype: int64
```

Explore class variable

```
df['y_Churn'].value_counts()

False.    3961
True.      656
Name: y_Churn, dtype: int64
```

The y_Churn is the target variable which is ordinal in nature.

```
ip_churn=df.groupby("c_InternationalPlan").y_Churn .value_counts()
ip_churn
```

```
c_InternationalPlan  y_Churn
no                  False.    3701
```

```

                True.      470
yes             False.    260
                True.     186
Name: y_Churn, dtype: int64

```

Correlation

```

correlation=df.corr().abs()
correlation

```

	c_AreaCode	q_AccountLength	q_CustServCalls	\
c_AreaCode	1.000000	0.020394	0.021046	
q_AccountLength	0.020394	1.000000	0.002620	
q_CustServCalls	0.021046	0.002620	1.000000	
q_DayCalls	0.013179	0.032783	0.008747	
q_DayCharge	0.018903	0.001999	0.008155	
q_DayMins	0.018900	0.002002	0.008149	
q_EveCalls	0.011528	0.015598	0.007730	
q_EveCharge	0.011533	0.006775	0.015611	
q_EveMins	0.011513	0.006778	0.015598	
q_InternationalCharge	0.007386	0.003501	0.016148	
q_InternationalMins	0.007292	0.003483	0.016079	
q_Internationalcalls	0.011531	0.023485	0.016778	
q_NightCalls	0.015316	0.009482	0.010258	
q_NightCharge	0.002782	0.002095	0.013868	
q_NightMins	0.002794	0.002077	0.013871	
q_VMailMessage	0.002597	0.012983	0.006951	
DC/min	0.014331	0.001995	0.018327	
EC/min	0.023385	0.014497	0.014045	
NC/min	0.004466	0.020796	0.001603	
IC/min	0.014661	0.011580	0.029188	

\	q_DayCalls	q_DayCharge	q_DayMins	q_EveCalls
c_AreaCode	0.013179	0.018903	0.018900	0.011528
q_AccountLength	0.032783	0.001999	0.002002	0.015598
q_CustServCalls	0.008747	0.008155	0.008149	0.007730
q_DayCalls	1.000000	0.002821	0.002823	0.003923
q_DayCharge	0.002821	1.000000	1.000000	0.012992
q_DayMins	0.002823	1.000000	1.000000	0.012990
q_EveCalls	0.003923	0.012992	0.012990	1.000000
q_EveCharge	0.006429	0.010262	0.010268	0.001151
q_EveMins	0.006430	0.010255	0.010260	0.001135

q_InternationalCharge	0.013055	0.012262	0.012261	0.002831
q_InternationalMins	0.012951	0.012315	0.012314	0.002798
q_Internationalcalls	0.010889	0.000163	0.000166	0.005198
q_NightCalls	0.013299	0.005164	0.005165	0.015463
q_NightCharge	0.010724	0.009593	0.009591	0.002624
q_NightMins	0.010730	0.009606	0.009604	0.002610
q_VMailMessage	0.003846	0.009025	0.009028	0.006508
DC/min	0.005004	0.045677	0.045904	0.005130
EC/min	0.009264	0.011786	0.011783	0.022399
NC/min	0.006709	0.013106	0.013099	0.010932
IC/min	0.008493	0.015359	0.015362	0.000118

q_InternationalCharge	q_EveCharge	q_EveMins	
c_AreaCode	\		
	0.011533	0.011513	0.007386
q_AccountLength	0.006775	0.006778	0.003501
q_CustServCalls	0.015611	0.015598	0.016148
q_DayCalls	0.006429	0.006430	0.013055
q_DayCharge	0.010262	0.010255	0.012262
q_DayMins	0.010268	0.010260	0.012261
q_EveCalls	0.001151	0.001135	0.002831
q_EveCharge	1.000000	1.000000	0.000170
q_EveMins	1.000000	1.000000	0.000172
q_InternationalCharge	0.000170	0.000172	1.000000
q_InternationalMins	0.000163	0.000165	0.999993

q_Internationalcalls	0.005751	0.005749	0.021548
q_NightCalls	0.014231	0.014226	0.006439
q_NightCharge	0.018846	0.018837	0.004497
q_NightMins	0.018839	0.018830	0.004500
q_VMailMessage	0.017878	0.017872	0.005202
DC/min	0.017246	0.017242	0.008715
EC/min	0.035968	0.036596	0.006301
NC/min	0.000705	0.000701	0.008879
IC/min	0.000293	0.000285	0.087671

	q_InternationalMins	q_Internationalcalls	\
c_AreaCode	0.007292	0.011531	
q_AccountLength	0.003483	0.023485	
q_CustServCalls	0.016079	0.016778	
q_DayCalls	0.012951	0.010889	
q_DayCharge	0.012315	0.000163	
q_DayMins	0.012314	0.000166	
q_EveCalls	0.002798	0.005198	
q_EveCharge	0.000163	0.005751	
q_EveMins	0.000165	0.005749	
q_InternationalCharge	0.999993	0.021548	
q_InternationalMins	1.000000	0.021431	
q_Internationalcalls	0.021431	1.000000	
q_NightCalls	0.006431	0.003294	
q_NightCharge	0.004546	0.014624	
q_NightMins	0.004549	0.014651	
q_VMailMessage	0.005167	0.007423	
DC/min	0.008690	0.011105	
EC/min	0.006301	0.005035	
NC/min	0.008852	0.026916	
IC/min	0.091300	0.027672	

	q_NightCalls	q_NightCharge	q_NightMins	\
c_AreaCode	0.015316	0.002782	0.002794	
q_AccountLength	0.009482	0.002095	0.002077	
q_CustServCalls	0.010258	0.013868	0.013871	
q_DayCalls	0.013299	0.010724	0.010730	
q_DayCharge	0.005164	0.009593	0.009606	
q_DayMins	0.005165	0.009591	0.009604	
q_EveCalls	0.015463	0.002624	0.002610	

q_EveCharge	0.014231	0.018846	0.018839
q_EveMins	0.014226	0.018837	0.018830
q_InternationalCharge	0.006439	0.004497	0.004500
q_InternationalMins	0.006431	0.004546	0.004549
q_Internationalcalls	0.003294	0.014624	0.014651
q_NightCalls	1.000000	0.025722	0.025742
q_NightCharge	0.025722	1.000000	0.999999
q_NightMins	0.025742	0.999999	1.000000
q_VMailMessage	0.000889	0.004665	0.004672
DC/min	0.028335	0.001849	0.001825
EC/min	0.007975	0.017241	0.017251
NC/min	0.004684	0.004907	0.003729
IC/min	0.001188	0.013655	0.013657

	q_VMailMessage	DC/min	EC/min	NC/min
IC/min				
c_AreaCode	0.002597	0.014331	0.023385	0.004466
0.014661				
q_AccountLength	0.012983	0.001995	0.014497	0.020796
0.011580				
q_CustServCalls	0.006951	0.018327	0.014045	0.001603
0.029188				
q_DayCalls	0.003846	0.005004	0.009264	0.006709
0.008493				
q_DayCharge	0.009025	0.045677	0.011786	0.013106
0.015359				
q_DayMins	0.009028	0.045904	0.011783	0.013099
0.015362				
q_EveCalls	0.006508	0.005130	0.022399	0.010932
0.000118				
q_EveCharge	0.017878	0.017246	0.035968	0.000705
0.000293				
q_EveMins	0.017872	0.017242	0.036596	0.000701
0.000285				
q_InternationalCharge	0.005202	0.008715	0.006301	0.008879
0.087671				
q_InternationalMins	0.005167	0.008690	0.006301	0.008852
0.091300				
q_Internationalcalls	0.007423	0.011105	0.005035	0.026916
0.027672				
q_NightCalls	0.000889	0.028335	0.007975	0.004684
0.001188				
q_NightCharge	0.004665	0.001849	0.017241	0.004907
0.013655				
q_NightMins	0.004672	0.001825	0.017251	0.003729
0.013657				
q_VMailMessage	1.000000	0.026199	0.010810	0.006164
0.006811				
DC/min	0.026199	1.000000	0.000187	0.020774
0.005734				

EC/min	0.010810	0.000187	1.000000	0.000557
0.010020				
NC/min	0.006164	0.020774	0.000557	1.000000
0.010549				
IC/min	0.006811	0.005734	0.010020	0.010549
1.000000				

#To find highly correlated continuos columns (Pearson Correlation is used for cont-cont columns)

```
hcorr_features=set()
```

```
x=0.95
```

```
for i in range(len(correlation.columns)):
```

```
    for j in range(i):
```

```
        if(correlation.iloc[i,j])>x:
```

```
            colname=correlation.columns[i]
```

```
            hcorr_features.add(colname)
```

```
print("Highly co-related continuos columns",hcorr_features)
```

Highly co-related continuos columns {'q_DayMins', 'q_EveMins', 'q_InternationalMins', 'q_NightMins'}

Dropping the highly correleated values from visulaization

```
df.drop('q_NightMins',axis=1,inplace=True)
```

```
df.drop('q_EveMins',axis=1,inplace=True)
```

```
df.drop('q_DayMins',axis=1,inplace=True)
```

```
df.drop('q_InternationalMins',axis=1,inplace=True)
```

df.drop('c_VMailPlan',axis=1,inplace=True) #This can be dropped because there are charges for subscribed customer and zero for unsubscribed customer

```
df.drop('c_AreaCode',axis=1,inplace=True)
```

```
df.drop('c_Phone',axis=1,inplace=True)
```

```
df.drop('q_AccountLength',axis=1,inplace=True)
```

```
df.drop('c_State',axis=1,inplace=True)
```

```
df.drop('q_CustServCalls',axis=1,inplace=True)
```

```
df.drop('DC/min',axis=1,inplace=True)
```

```
df.drop('EC/min',axis=1,inplace=True)
```

```
df.drop('NC/min',axis=1,inplace=True)
```

```
df.drop('IC/min',axis=1,inplace=True)
```

```
df
```

	c_InternationalPlan	q_DayCalls	q_DayCharge	q_EveCalls
q_EveCharge \				
0	no	110	45.07	99
16.78				
1	no	123	27.47	103
16.62				
2	no	114	41.38	110
10.30				
3	yes	71	50.90	88
5.26				

4	yes	113	28.34	122
12.61				
...
...				
4612	no	81	24.48	112
15.91				
4613	no	91	32.13	96
25.76				
4614	no	126	21.68	129
25.17				
4615	no	98	28.71	117
19.24				
4616	no	100	34.80	107
10.78				

	q_InternationalCharge	q_Internationalcalls	q_NightCalls	\
0	2.70	3	91	
1	3.70	3	103	
2	3.29	5	104	
3	1.78	7	89	
4	2.73	3	121	
...	
4612	2.30	6	122	
4613	4.24	1	116	
4614	3.51	3	91	
4615	3.86	3	96	
4616	3.27	4	115	

	q_NightCharge	q_VMailMessage	y_Churn
0	11.01	25	False.
1	11.45	26	False.
2	7.32	0	False.
3	8.86	0	False.
4	8.41	0	False.
...
4612	7.14	25	False.
4613	7.36	29	False.
4614	9.04	33	False.
4615	7.45	23	False.
4616	9.13	0	False.

[4617 rows x 11 columns]

Encoding

```
from sklearn.preprocessing import LabelEncoder
l_enc = LabelEncoder()
for i in (0,1,1):
    df.iloc[:,i] = l_enc.fit_transform(df.iloc[:,i])
df
```


q_EveCharge \	c_InternationalPlan	q_DayCalls	q_DayCharge	q_EveCalls
0	0	74	45.07	99
16.78				
1	0	87	27.47	103
16.62				
2	0	78	41.38	110
10.30				
3	1	35	50.90	88
5.26				
4	1	77	28.34	122
12.61				
...
...				
4612	0	45	24.48	112
15.91				
4613	0	55	32.13	96
25.76				
4614	0	90	21.68	129
25.17				
4615	0	62	28.71	117
19.24				
4616	0	64	34.80	107
10.78				

	q_InternationalCharge	q_Internationalcalls	q_NightCalls \
0	2.70	3	91
1	3.70	3	103
2	3.29	5	104
3	1.78	7	89
4	2.73	3	121
...
4612	2.30	6	122
4613	4.24	1	116
4614	3.51	3	91
4615	3.86	3	96
4616	3.27	4	115

	q_NightCharge	q_VMailMessage	y_Churn
0	11.01	25	False.
1	11.45	26	False.
2	7.32	0	False.
3	8.86	0	False.
4	8.41	0	False.
...
4612	7.14	25	False.
4613	7.36	29	False.
4614	9.04	33	False.
4615	7.45	23	False.
4616	9.13	0	False.

```
[4617 rows x 11 columns]
```

```
y = df.y_Churn  
X = df.iloc[:, :-1]
```

Splitting datas into train and test datas

```
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test =  
train_test_split(X,y,test_size=0.2,random_state=1)
```

Feature engineering

- checking the data types of the variable after splitting

```
# check data types in X_train
```

```
X_train.dtypes
```

```
c_InternationalPlan      int32  
q_DayCalls               int64  
q_DayCharge              float64  
q_EveCalls               int64  
q_EveCharge              float64  
q_InternationalCharge    float64  
q_Internationalcalls     int64  
q_NightCalls             int64  
q_NightCharge            float64  
q_VMailMessage           int64  
dtype: object
```

Train the model using DecisionTreeClassifier

```
from sklearn.tree import DecisionTreeClassifier  
clf = DecisionTreeClassifier()  
clf = clf.fit(X_train,y_train)  
y_pred = clf.predict(X_test)
```

Evaluating the accuracy of the model

```
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

```
Accuracy: 0.8560606060606061
```

Optimizing the performance of the decision tree model

- Criteria choosen is "entropy" and "gini"
- max_depth is given as 3 for our model

```
Decision Tree Classifier with criterion entropy index
```

```
# Create Decision Tree classifer object
```

```
clf_en = DecisionTreeClassifier(criterion="entropy", max_depth=3)
```

```
# Train Decision Tree Classifier
```

```

clf_en = clf_en.fit(X_train,y_train)

#Predict the response for test dataset
y_pred = clf_en.predict(X_test)

# Model Accuracy, how often is the classifier correct?
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))

Accuracy: 0.9036796536796536

```

Here, y_test are the true class labels and y_pred are the predicted class labels in the test-set.

Compare the train-set and test-set accuracy

Now, we are comparing the train-set and test-set accuracy to check for overfitting.

```

y_pred_train_entropy = clf_en.predict(X_train)

y_pred_train_entropy
array([' False.', ' False.', ' False.', ..., ' False.', ' False.',
       ' False.'], dtype=object)

from sklearn.metrics import accuracy_score
print('Training-set accuracy score: {0:0.4f}'.
      format(accuracy_score(y_train, y_pred_train_entropy)))

Training-set accuracy score: 0.8993

```

Check for overfitting and underfitting

```

print('Training set score: {:.4f}'.format(clf_en.score(X_train,
y_train)))

print('Test set score: {:.4f}'.format(clf_en.score(X_test, y_test)))

Training set score: 0.8993
Test set score: 0.9037

```

Here, the training-set accuracy score is 0.8985 while the test-set accuracy is 0.9048. These two values are quite comparable. So, there is no sign of overfitting.

Decision Tree Classifier with criterion gini index

```

# Create Decision Tree classifier object
clf = DecisionTreeClassifier(criterion="gini", max_depth=3)

# Train Decision Tree Classifier
clf = clf.fit(X_train,y_train)

#Predict the response for test dataset
y_pred = clf.predict(X_test)

```

```
# Model Accuracy, how often is the classifier correct?  
print("Accuracy:", metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.9047619047619048

Compare the train-set and test-set accuracy

Now, we are comparing the train-set and test-set accuracy to check for overfitting.

```
y_pred_train_gini = clf.predict(X_train)  
  
y_pred_train_gini  
array([' False.', ' False.', ' False.', ..., ' False.', ' False.',  
       ' False.'], dtype=object)
```

```
from sklearn.metrics import accuracy_score  
print('Training-set accuracy score: {0:0.4f}'.  
      format(accuracy_score(y_train, y_pred_train_gini)))
```

Training-set accuracy score: 0.8985

Check for overfitting and underfitting

```
print('Training set score: {:.4f}'.format(clf.score(X_train,  
y_train)))  
  
print('Test set score: {:.4f}'.format(clf.score(X_test, y_test)))
```

Training set score: 0.8985
Test set score: 0.9048

We can see that the training-set score and test-set score is same as above. The training-set accuracy score is 0.8985 while the test-set accuracy to be 0.9048. These two values are quite comparable. So, there is no sign of overfitting.

Now, based on the above analysis we can conclude that our classification model accuracy is very good.

But, it does not give the underlying distribution of values. Also, it does not tell anything about the type of errors our classifier is making.

We have another tool called Confusion matrix that comes to our rescue.

Summary:

-In this project, we have build a Decision-Tree Classifier model to predict the churn rate. - We build two models, one with criterion gini index and another one with criterion entropy. The model yields a very good performance as indicated by the model accuracy in both the cases which was found to be 0.90. -In the model with criterion gini index, the training-set accuracy score is 0.8985 while the test-set accuracy to be 0.9048. These two values are quite comparable. So, there is no sign of overfitting. -Similarly, in the model with criterion

entropy, the training-set accuracy score is 0.8993 while the test-set accuracy to be 0.9037. We get the same values as in the case with criterion gini. So, there is no sign of overfitting.